

R Notebook

[http://ijlyttle.github.io/isugg_purrr/presentation.html#\(1\)](http://ijlyttle.github.io/isugg_purrr/presentation.html#(1))

Packages to run this presentation

```
library("readr")
library("tibble")
library("dplyr")
library("tidyr")
library("stringr")
library("ggplot2")
library("purrr")
library("broom")
```

Motivation

As you know, purrr is a recent package from Hadley Wickham, focused on lists and functional programming, like dplyr is focused on data-frames.

I figure a good way to learn a new package is to try to solve a problem, so we have a dataset:

- you can view or download
- you can download the source of this presentation
- these are three temperatures recorded simultaneously in a piece of electronics
- it will be very valuable to be able to characterize the transient temperature for each sensor
- we want to apply the same set of models across all three sensors
- it will be easier to show using pictures

Let's get the data into shape

Using the readr package

```
temperature_wide <-
  read_csv("temperature.csv") %>%
  print()
```

```
# A tibble: 327 x 4
  instant      temperature_a temperature_b temperature_c
  <dtm>          <dbl>          <dbl>          <dbl>
1 2015-11-13 06:10:19      116.           91.7           84.2
2 2015-11-13 06:10:23      116.           91.7           84.2
3 2015-11-13 06:10:27      116.           91.6           84.2
4 2015-11-13 06:10:31      116.           91.7           84.2
5 2015-11-13 06:10:36      116.           91.7           84.2
6 2015-11-13 06:10:41      116.           91.6           84.2
7 2015-11-13 06:10:46      116.           91.5           84.2
8 2015-11-13 06:10:51      116.           91.5           84.2
```

```

  9 2015-11-13 06:10:56      116.      91.5      84.2
10 2015-11-13 06:11:01      115.      91.5      84.2
# ... with 317 more rows

```

Is temperature_wide “tidy”?

```

# A tibble: 327 x 4
  instant      temperature_a temperature_b temperature_c
  <dtm>          <dbl>          <dbl>          <dbl>
1 2015-11-13 06:10:19      116.      91.7      84.2
2 2015-11-13 06:10:23      116.      91.7      84.2
3 2015-11-13 06:10:27      116.      91.6      84.2
4 2015-11-13 06:10:31      116.      91.7      84.2
5 2015-11-13 06:10:36      116.      91.7      84.2
6 2015-11-13 06:10:41      116.      91.6      84.2
7 2015-11-13 06:10:46      116.      91.5      84.2
8 2015-11-13 06:10:51      116.      91.5      84.2
9 2015-11-13 06:10:56      116.      91.5      84.2
10 2015-11-13 06:11:01      115.      91.5      84.2
# ... with 317 more rows

```

Why or why not?

Tidy data

1. Each column is a variable
2. Each row is an observation
3. Each cell is a value

(<http://www.jstatsoft.org/v59/i10/paper>)

My personal observation is that “tidy” can depend on the context, on what you want to do with the data.

Let’s get this into a tidy form

```

temperature_tall <-
  temperature_wide %>%
  gather(key = "id_sensor", value = "temperature", starts_with("temp")) %>%
  mutate(id_sensor = str_replace(id_sensor, "temperature_", "")) %>%
  print()

```

```

# A tibble: 981 x 3
  instant      id_sensor temperature
  <dtm>          <chr>          <dbl>
1 2015-11-13 06:10:19 a             116.
2 2015-11-13 06:10:23 a             116.
3 2015-11-13 06:10:27 a             116.
4 2015-11-13 06:10:31 a             116.
5 2015-11-13 06:10:36 a             116.
6 2015-11-13 06:10:41 a             116.
7 2015-11-13 06:10:46 a             116.
8 2015-11-13 06:10:51 a             116.
9 2015-11-13 06:10:56 a             116.


```

```
10 2015-11-13 06:11:01 a
# ... with 971 more rows
```

115.

Now, it's easier to visualize

```
temperature_tall %>%
  ggplot(aes(x = instant, y = temperature, color = id_sensor)) +
  geom_line()
```



Rearrange a bit more

`delta_time` Δt

change in time since event started, s

`delta_temperature`: ΔT


change in temperature since event started, °C

```
delta <-
  temperature_tall %>%
  arrange(id_sensor, instant) %>%
  group_by(id_sensor) %>%
  mutate(
    delta_time = as.numeric(instant) - as.numeric(instant[[1]]),
    delta_temperature = temperature - temperature[[1]]
```

```
) %>%
select(id_sensor, delta_time, delta_temperature)
```

Let's have a look

```
delta %>%
  ggplot(aes(x = delta_time, y = delta_temperature, color = id_sensor)) +
  geom_line()
```



Curve-fitting

We want to see how three different curve-fits might perform on these three data-sets:

Newtonian cooling

$$\Delta T = \Delta T_0$$

$$\Delta T = \Delta T_0 \left[1 - \exp \left(- \frac{\Delta T}{\Delta T_0} \right) \right]$$

Semi-infinite solid

$$\Delta T = \Delta T_0 \operatorname{erfc} \left(\sqrt{\frac{\Delta T}{\Delta T_0}} \right)$$

Semi-infinite solid with convection

$\Delta T = \Delta T_0 \left[\operatorname{erfc} \left(\sqrt{\frac{\tau_0}{\Delta t}} \right) - \exp \left(\frac{Bi_0}{2} \right) \right]$

Some definitions

```
# reference: http://stackoverflow.com/questions/29067916/r-error-function-erfz
# (see Abramowitz and Stegun 29.2.29)
erf <- function(x) 2 * pnorm(x * sqrt(2)) - 1
erfc <- function(x) 2 * pnorm(x * sqrt(2), lower = FALSE)
```

```
newton_cooling <- function(x) {
  nls(
    delta_temperature ~ delta_temperature_0 * (1 - exp(-delta_time/tau_0)),
    start = list(delta_temperature_0 = -10, tau_0 = 50),
    data = x
  )
}
```

More math

```
semi_infinite_simple <- function(x) {
  nls(
    delta_temperature ~ delta_temperature_0 * erfc(sqrt(tau_0/delta_time)),
    start = list(delta_temperature_0 = -10, tau_0 = 50),
    data = x
  )
}

semi_infinite_convection <- function(x){
  nls(
    delta_temperature ~
      delta_temperature_0 * (
        erfc(sqrt(tau_0/delta_time)) -
        exp(Bi_0 + (Bi_0/2)^2 * delta_time/tau_0) *
        erfc(sqrt(tau_0/delta_time) +
          (Bi_0/2) * sqrt(delta_time/tau_0))
      ),
    start = list(delta_temperature_0 = -5, tau_0 = 50, Bi_0 = 1.e6),
    data = x
  )
}
```

Before we get into purrr

Before doing anything, we want to show that we can do something with one dataset and one model-function:

```
tmp_data <- delta %>% filter(id_sensor == "a")

tmp_model <- newton_cooling(tmp_data)
```

```
summary(tmp_model)
```

Formula: $\text{delta_temperature} \sim \text{delta_temperature}_0 * (1 - \exp(-\text{delta_time}/\text{tau}_0))$

Parameters:

	Estimate	Std. Error	t value	Pr(> t)
delta_temperature_0	-15.06085	0.05262	-286.2	<2e-16 ***
tau_0	500.01382	4.83673	103.4	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3267 on 325 degrees of freedom

Number of iterations to convergence: 7

Achieved convergence tolerance: 4.136e-06


Look at predictions

```
tmp_pred <-  
  tmp_data %>%  
  mutate(modeled = predict(tmp_model, data = .)) %>%  
  select(id_sensor, delta_time, measured = delta_temperature, modeled) %>%  
  gather("type", "delta_temperature", measured:modeled) %>%  
  print()
```

```
# A tibble: 654 x 4  
# Groups:   id_sensor [1]  
  id_sensor delta_time type      delta_temperature  
  <chr>      <dbl> <chr>      <dbl>  
1 a          0 measured          0  
2 a          4 measured          0  
3 a          8 measured     -0.06  
4 a         12 measured     -0.06  
5 a         17 measured    -0.211  
6 a         22 measured    -0.423  
7 a         27 measured    -0.423  
8 a         32 measured    -0.574  
9 a         37 measured    -0.726  
10 a        42 measured    -0.878  
# ... with 644 more rows
```

A more-useful look

```
tmp_pred %>%  
  ggplot(aes(x = delta_time, y = delta_temperature, linetype = type)) +  
  geom_line()
```



“Regular” data-frame

```
print(delta)
```

```
# A tibble: 981 x 3
# Groups:   id_sensor [3]
  id_sensor delta_time delta_temperature
  <chr>      <dbl>      <dbl>
1 a          0          0
2 a          4          0
3 a          8         -0.06
4 a         12         -0.06
5 a         17        -0.211
6 a         22        -0.423
7 a         27        -0.423
8 a         32        -0.574
9 a         37        -0.726
10 a        42        -0.878
# ... with 971 more rows
```

Each column of the dataframe is a vector - in this case, a character vector and two doubles

How to make a weird data-frame

Here's where the fun starts - a column of a data-frame can be a list.

- use `tidyr::nest()` to makes a column `data`, which is a list of data-frames
- this seems like a stronger expression of the `dplyr::group_by()` idea

```
delta_nested <-  
  delta %>%  
  nest(-id_sensor) %>%  
  print()
```

```
# A tibble: 3 x 2  
  id_sensor data  
  <chr>     <list>  
1 a       <tibble [327 x 2]>  
2 b       <tibble [327 x 2]>  
3 c       <tibble [327 x 2]>
```

Map data-frames to the modeling function

- `map()` is like `lapply()`
- `map()` returns a list-column (it keeps the weirdness)

```
model_nested <-  
  delta_nested %>%  
  mutate(model = map(data, newton_cooling)) %>%  
  print()
```

```
# A tibble: 3 x 3  
  id_sensor data          model  
  <chr>     <list>         <list>  
1 a       <tibble [327 x 2]> <S3: nls>  
2 b       <tibble [327 x 2]> <S3: nls>  
3 c       <tibble [327 x 2]> <S3: nls>
```

We can use `map2()` to make the predictions

- `map2()` is like `mapapply()`
- designed to map two columns (`model`, `data`) to a function `predict()`

```
predict_nested <-  
  model_nested %>%  
  mutate(pred = map2(model, data, predict)) %>%  
  print()
```

```
# A tibble: 3 x 4  
  id_sensor data          model    pred  
  <chr>     <list>         <list>  <list>  
1 a       <tibble [327 x 2]> <S3: nls> <dbl [327]>  
2 b       <tibble [327 x 2]> <S3: nls> <dbl [327]>  
3 c       <tibble [327 x 2]> <S3: nls> <dbl [327]>
```


We need to get out of the weirdness

- use `unnest()` to get back to a regular data-frame

```
predict_unnested <-  
  predict_nested %>%  
  unnest(data, pred) %>%  
  print()  
  
# A tibble: 981 x 4  
  id_sensor    pred delta_time delta_temperature  
  <chr>      <dbl>    <dbl>          <dbl>  
1 a          0         0            0  
2 a        -0.120         4            0  
3 a        -0.239         8        -0.06  
4 a        -0.357        12        -0.06  
5 a        -0.503        17       -0.211  
6 a        -0.648        22       -0.423  
7 a        -0.792        27       -0.423  
8 a        -0.934        32       -0.574  
9 a         -1.07        37       -0.726  
10 a        -1.21        42       -0.878  
# ... with 971 more rows
```

We can wrangle the predictions


- get into a form that makes it easier to plot

```
predict_tall <-  
  predict_unnested %>%  
  rename(modeled = pred, measured = delta_temperature) %>%  
  gather("type", "delta_temperature", modeled, measured) %>%  
  print()  
  
# A tibble: 1,962 x 4  
  id_sensor delta_time type    delta_temperature  
  <chr>      <dbl> <chr>          <dbl>  
1 a          0 modeled            0  
2 a          4 modeled       -0.120  
3 a          8 modeled       -0.239  
4 a         12 modeled       -0.357  
5 a         17 modeled       -0.503  
6 a         22 modeled       -0.648  
7 a         27 modeled       -0.792  
8 a         32 modeled       -0.934  
9 a         37 modeled        -1.07  
10 a         42 modeled        -1.21  
# ... with 1,952 more rows
```

We can visualize the predictions

```
predict_tall %>%  
  ggplot(aes(x = delta_time, y = delta_temperature)) +
```

```
geom_line(aes(color = id_sensor, linetype = type))
```



Now we want to look at a selection of models

Make a list of functions to model:

```
list_model <-
  list(
    newton_cooling = newton_cooling,
    semi_infinite_simple = semi_infinite_simple,
    semi_infinite_convection = semi_infinite_convection
  )
```

Step: write a function to define the “inner” loop

```
fn_model <- function(.model, df){
  # safer to avoid non-standard evaluation
  # df %>% mutate(model = map(data, .model))

  df$model <- map(df$data, possibly(.model, NULL))
  df
}
```

- for a given model-function and a given (weird) data-frame, return a modified version of that data-frame

with a column `model`, which is the model-function applied to each element of the data-frame's `data` column (which is itself a list of data-frames)

- the purrr functions `safely()` and `possibly()` are **very** interesting. I think they could be useful outside of purrr as a friendlier way to do error-handling.

Step: `map_df()` to define the “outer” loop

```
model_nested_new <-
  list_model %>%
  map_df(fn_model, delta_nested, .id = "id_model") %>%
  print()
```

```
# A tibble: 9 x 4
  id_model      id_sensor data              model
  <chr>        <chr>    <list>          <list>
1 newton_cooling a      <tibble [327 x 2]> <S3: nls>
2 newton_cooling b      <tibble [327 x 2]> <S3: nls>
3 newton_cooling c      <tibble [327 x 2]> <S3: nls>
4 semi_infinite_simple a    <tibble [327 x 2]> <S3: nls>
5 semi_infinite_simple b    <tibble [327 x 2]> <S3: nls>
6 semi_infinite_simple c    <tibble [327 x 2]> <S3: nls>
7 semi_infinite_convection a  <tibble [327 x 2]> <NULL>
8 semi_infinite_convection b  <tibble [327 x 2]> <NULL>
9 semi_infinite_convection c  <tibble [327 x 2]> <NULL>
```

- for each element of a list of model-functions, run the inner-loop function, and row-bind the results into a data-frame
- we want to discard the rows where the model failed
- we also want to investigate why they failed, but that's a different talk

Step: `map()` to identify the null models

```
model_nested_new <-
  list_model %>%
  map_df(fn_model, delta_nested, .id = "id_model") %>%
  mutate(is_null = map(model, is.null)) %>%
  print()
```

```
# A tibble: 9 x 5
  id_model      id_sensor data              model      is_null
  <chr>        <chr>    <list>          <list>    <list>
1 newton_cooling a      <tibble [327 x 2]> <S3: nls> <lgl [1]>
2 newton_cooling b      <tibble [327 x 2]> <S3: nls> <lgl [1]>
3 newton_cooling c      <tibble [327 x 2]> <S3: nls> <lgl [1]>
4 semi_infinite_simple a    <tibble [327 x 2]> <S3: nls> <lgl [1]>
5 semi_infinite_simple b    <tibble [327 x 2]> <S3: nls> <lgl [1]>
6 semi_infinite_simple c    <tibble [327 x 2]> <S3: nls> <lgl [1]>
7 semi_infinite_convection a  <tibble [327 x 2]> <NULL> <lgl [1]>
8 semi_infinite_convection b  <tibble [327 x 2]> <NULL> <lgl [1]>
9 semi_infinite_convection c  <tibble [327 x 2]> <NULL> <lgl [1]>
```

- using `map(model, is.null)` returns a list column
- to use `filter()`, we have to escape the weirdness

Step: `map_lgl()` to identify nulls and get out of the weirdness

```
model_nested_new <-
  list_model %>%
  map_df(fn_model, delta_nested, .id = "id_model") %>%
  mutate(is_null = map_lgl(model, is.null)) %>%
  print()
```

A tibble: 9 x 5

	id_model <chr>	id_sensor <chr>	data <list>	model <list>	is_null <lgl>
1	newton_cooling	a	<tibble [327 x 2]>	<S3: nls>	FALSE
2	newton_cooling	b	<tibble [327 x 2]>	<S3: nls>	FALSE
3	newton_cooling	c	<tibble [327 x 2]>	<S3: nls>	FALSE
4	semi_infinite_simple	a	<tibble [327 x 2]>	<S3: nls>	FALSE
5	semi_infinite_simple	b	<tibble [327 x 2]>	<S3: nls>	FALSE
6	semi_infinite_simple	c	<tibble [327 x 2]>	<S3: nls>	FALSE
7	semi_infinite_convection	a	<tibble [327 x 2]>	<NULL>	TRUE
8	semi_infinite_convection	b	<tibble [327 x 2]>	<NULL>	TRUE
9	semi_infinite_convection	c	<tibble [327 x 2]>	<NULL>	TRUE

- using `map_lgl(model, is.null)` returns a vector column

Step: `filter()` and `select()` to clean up

```
model_nested_new <-
  list_model %>%
  map_df(fn_model, delta_nested, .id = "id_model") %>%
  mutate(is_null = map_lgl(model, is.null)) %>%
  filter(!is_null) %>%
  select(-is_null) %>%
  print()
```

A tibble: 6 x 4

	id_model <chr>	id_sensor <chr>	data <list>	model <list>
1	newton_cooling	a	<tibble [327 x 2]>	<S3: nls>
2	newton_cooling	b	<tibble [327 x 2]>	<S3: nls>
3	newton_cooling	c	<tibble [327 x 2]>	<S3: nls>
4	semi_infinite_simple	a	<tibble [327 x 2]>	<S3: nls>
5	semi_infinite_simple	b	<tibble [327 x 2]>	<S3: nls>
6	semi_infinite_simple	c	<tibble [327 x 2]>	<S3: nls>

Let's get predictions

```
predict_nested <-
  model_nested_new %>%
```

```
mutate(pred = map2(model, data, predict)) %>%
print()
```

```
# A tibble: 6 x 5
  id_model      id_sensor data          model      pred
  <chr>         <chr>    <list>      <list>    <list>
1 newton_cooling a      <tibble [327 x 2]> <S3: nls> <dbl [327]>
2 newton_cooling b      <tibble [327 x 2]> <S3: nls> <dbl [327]>
3 newton_cooling c      <tibble [327 x 2]> <S3: nls> <dbl [327]>
4 semi_infinite_simple a    <tibble [327 x 2]> <S3: nls> <dbl [327]>
5 semi_infinite_simple b    <tibble [327 x 2]> <S3: nls> <dbl [327]>
6 semi_infinite_simple c    <tibble [327 x 2]> <S3: nls> <dbl [327]>
```


unnest(), make it tall

```
predict_tall <-
  predict_nested %>%
  unnest(data, pred) %>%
  rename(modeled = pred, measured = delta_temperature) %>%
  gather("type", "delta_temperature", modeled, measured) %>%
  print()
```

```
# A tibble: 3,924 x 5
  id_model      id_sensor delta_time type      delta_temperature
  <chr>         <chr>      <dbl> <chr>      <dbl>
1 newton_cooling a          0 modeled          0
2 newton_cooling a          4 modeled        -0.120
3 newton_cooling a          8 modeled        -0.239
4 newton_cooling a         12 modeled        -0.357
5 newton_cooling a         17 modeled        -0.503
6 newton_cooling a         22 modeled        -0.648
7 newton_cooling a         27 modeled        -0.792
8 newton_cooling a         32 modeled        -0.934
9 newton_cooling a         37 modeled        -1.07
10 newton_cooling a         42 modeled        -1.21
# ... with 3,914 more rows
```

We can visualize the predictions

```
predict_tall %>%
  ggplot(aes(x = delta_time, y = delta_temperature)) +
  geom_line(aes(color = id_sensor, linetype = type)) +
  facet_grid(id_model ~ .)
```




Let's get the residuals

```
resid <-
  model_nested_new %>%
  mutate(resid = map(model, resid)) %>%
  unnest(data, resid) %>%
  print()
```

```
# A tibble: 1,962 x 5
  id_model id_sensor resid delta_time delta_temperature
  <chr>    <chr>    <dbl>    <dbl>    <dbl>
1 newton_cooling a      0         0         0
2 newton_cooling a    0.120         4         0
3 newton_cooling a    0.179         8    -0.06
4 newton_cooling a    0.297        12    -0.06
5 newton_cooling a    0.292        17   -0.211
6 newton_cooling a    0.225        22   -0.423
7 newton_cooling a    0.369        27   -0.423
8 newton_cooling a    0.360        32   -0.574
9 newton_cooling a    0.348        37   -0.726
10 newton_cooling a    0.335        42   -0.878
# ... with 1,952 more rows
```

And visualize them

```
resid %>%
  ggplot(aes(x = delta_time, y = resid)) +
  geom_line(aes(color = id_sensor)) +
  facet_grid(id_model ~ .)
```



Using broom package to look at model-statistics

The tidy() function extracts statistics from a model

```
model_parameters <-
  model_nested_new %>%
  select(id_model, id_sensor, model) %>%
  mutate(tidy = map(model, tidy)) %>%
  select(-model) %>%
  unnest() %>%
  print()
```

```
# A tibble: 12 x 7
  id_model id_sensor term estimate std.error statistic p.value
  <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl>
1 newton_cool~ a delta_te~ -15.1 0.0526 -286. 0.
2 newton_cool~ a tau_0 500. 4.84 103. 1.07e-250
3 newton_cool~ b delta_te~ -7.59 0.0676 -112. 6.38e-262
```

4	newton_cool~ b	tau_0	1041.	16.2	64.2	9.05e-187
5	newton_cool~ c	delta_te~	-9.87	0.704	-14.0	3.16e- 35
6	newton_cool~ c	tau_0	3525.	299.	11.8	5.61e- 27
7	semi_infini~ a	delta_te~	-21.5	0.0649	-332.	0.
8	semi_infini~ a	tau_0	139.	1.15	121.	2.14e-272
9	semi_infini~ b	delta_te~	-10.6	0.0515	-206.	0.
10	semi_infini~ b	tau_0	287.	2.58	111.	1.46e-260
11	semi_infini~ c	delta_te~	-8.04	0.0129	-626.	0.
12	semi_infini~ c	tau_0	500.	1.07	468.	0.

Get a sense of the coefficients

```
model_summary <-
  model_parameters %>%
  select(id_model, id_sensor, term, estimate) %>%
  spread(key = "term", value = "estimate") %>%
  print()
```

```
# A tibble: 6 x 4
  id_model      id_sensor delta_temperature_0 tau_0
  <chr>         <chr>          <dbl> <dbl>
1 newton_cooling a             -15.1  500.
2 newton_cooling b              -7.59 1041.
3 newton_cooling c              -9.87 3525.
4 semi_infinite_simple a        -21.5  139.
5 semi_infinite_simple b        -10.6  287.
6 semi_infinite_simple c         -8.04  500.
```

Summary

- this is just a small part of purrr
- there seem to be parallels between `tidyr::nest()`/`purrr::map()` and `dplyr::group_by()`/`dplyr::do()`
 - to my mind, the purrr framework is more understandable
 - update tweet from Hadley

References from Hadley:

- purrr 0.1.0 announcement
- purrr 0.2.0 announcement
- chapter from Garrett Golemund and Hadley's forthcoming book